

▼ **Australia Rain Prediction**

```
from google.colab import drive
drive.mount("/content/drive")
```

Mounted at /content/drive

```
import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import seaborn as sns
import warnings
warnings.filterwarnings("ignore")
```

```
df=pd.read_csv("/content/drive/MyDrive/Dataset/weatherAUS.csv")
```

```
df.head()
```

	Date	Location	MinTemp	MaxTemp	Rainfall	Evaporation	Sunshine	WindGustDir	WindGustSpeed	WindC
0	2008-12-01	Albury	13.4	22.9	0.6	NaN	NaN	W	44.0	
1	2008-12-02	Albury	7.4	25.1	0.0	NaN	NaN	WNW	44.0	
2	2008-12-03	Albury	12.9	25.7	0.0	NaN	NaN	WSW	46.0	
3	2008-12-04	Albury	9.2	28.0	0.0	NaN	NaN	NE	24.0	
4	2008-12-05	Albury	17.5	32.3	1.0	NaN	NaN	W	41.0	

5 rows × 23 columns



```
df.isnull().sum()/len(df)*100
```

Date	0.000000
Location	0.000000
MinTemp	1.020899
MaxTemp	0.866905

```

Rainfall      2.241853
Evaporation   43.166506
Sunshine      48.009762
WindGustDir   7.098859
WindGustSpeed 7.055548
WindDir9am    7.263853
WindDir3pm    2.906641
WindSpeed9am  1.214767
WindSpeed3pm  2.105046
Humidity9am   1.824557
Humidity3pm   3.098446
Pressure9am   10.356799
Pressure3pm   10.331363
Cloud9am      38.421559
Cloud3pm      40.807095
Temp9am       1.214767
Temp3pm       2.481094
RainToday     2.241853
RainTomorrow  2.245978
dtype: float64

```

```
df["Location"].value_counts()
```

```

Canberra      3436
Sydney        3344
Darwin         3193
Melbourne     3193
Brisbane      3193
Adelaide      3193
Perth         3193
Hobart        3193
Albany        3040
MountGambier  3040
Ballarat      3040
Townsville    3040
GoldCoast     3040
Cairns        3040
Launceston    3040
AliceSprings  3040
Bendigo       3040
Albury        3040
MountGinini   3040
Wollongong    3040
Newcastle     3039
Tuggeranong   3039
Penrith       3039
Woomera       3009
Nuriootpa     3009
Cobar         3009
CoffsHarbour  3009
Moree         3009
Sale          3009
PerthAirport  3009
PearceRAAF    3009

```

```

Witchcliffe      3009
BadgerysCreek    3009
Mildura          3009
NorfolkIsland    3009
MelbourneAirport 3009
Richmond         3009
SydneyAirport    3009
Waggawagga       3009
Williamtown      3009
Dartmoor         3009
Watsonia         3009
Portland         3009
Walpole          3006
NorahHead        3004
SalmonGums       3001
Katherine        1578
Nhil             1578
Uluru            1578
Name: Location, dtype: int64

```

```
df["Location"].duplicated()
```

```

0      False
1       True
2       True
3       True
4       True
...
145455   True
145456   True
145457   True
145458   True
145459   True
Name: Location, Length: 145460, dtype: bool

```

```
n_cols=df.select_dtypes(include=['int','float']).columns
n_cols
```

```

Index(['MinTemp', 'MaxTemp', 'Rainfall', 'Evaporation', 'Sunshine',
      'WindGustSpeed', 'WindSpeed9am', 'WindSpeed3pm', 'Humidity9am',
      'Humidity3pm', 'Pressure9am', 'Pressure3pm', 'Cloud9am', 'Cloud3pm',
      'Temp9am', 'Temp3pm'],
      dtype='object')

```

```
c_cols=df.select_dtypes(include=['object']).columns
c_cols
```

```

Index(['Date', 'Location', 'WindGustDir', 'WindDir9am', 'WindDir3pm',
      'RainToday', 'RainTomorrow'],
      dtype='object')

```

```

dtype='object')

for i in c_cols:
    print(i, df[i].isnull().sum())

    Date 0
    Location 0
    WindGustDir 10326
    WindDir9am 10566
    WindDir3pm 4228
    RainToday 3261
    RainTomorrow 3267

for i in c_cols:
    df[i].fillna(df[i].mode()[0], inplace=True)

for i in n_cols:
    print(i, df[i].isnull().sum())

    MinTemp 1485
    MaxTemp 1261
    Rainfall 3261
    Evaporation 62790
    Sunshine 69835
    WindGustSpeed 10263
    WindSpeed9am 1767
    WindSpeed3pm 3062
    Humidity9am 2654
    Humidity3pm 4507
    Pressure9am 15065
    Pressure3pm 15028
    Cloud9am 55888
    Cloud3pm 59358
    Temp9am 1767
    Temp3pm 3609

for i in n_cols:
    df[i].fillna(df[i].median(), inplace=True)

df.isnull().sum()

    Date      0
    Location  0
    MinTemp   0
    MaxTemp   0
    Rainfall   0
    Evaporation 0
    Sunshine   0
    WindGustDir 0
    WindGustSpeed 0

```

```

WindDir9am      0
WindDir3pm      0
WindSpeed9am    0
WindSpeed3pm    0
Humidity9am     0
Humidity3pm     0
Pressure9am     0
Pressure3pm     0
Cloud9am        0
Cloud3pm        0
Temp9am         0
Temp3pm         0
RainToday       0
RainTomorrow    0
dtype: int64

```

```
df.info()
```


```

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 145460 entries, 0 to 145459
Data columns (total 23 columns):
 #   Column              Non-Null Count  Dtype
---  -
 0   Date                145460 non-null object
 1   Location            145460 non-null object
 2   MinTemp             145460 non-null float64
 3   MaxTemp             145460 non-null float64
 4   Rainfall            145460 non-null float64
 5   Evaporation         145460 non-null float64
 6   Sunshine            145460 non-null float64
 7   WindGustDir         145460 non-null object
 8   WindGustSpeed       145460 non-null float64
 9   WindDir9am         145460 non-null object
10   WindDir3pm         145460 non-null object
11   WindSpeed9am       145460 non-null float64
12   WindSpeed3pm       145460 non-null float64
13   Humidity9am        145460 non-null float64
14   Humidity3pm        145460 non-null float64
15   Pressure9am        145460 non-null float64
16   Pressure3pm        145460 non-null float64
17   Cloud9am           145460 non-null float64
18   Cloud3pm           145460 non-null float64
19   Temp9am            145460 non-null float64
20   Temp3pm            145460 non-null float64
21   RainToday          145460 non-null object
22   RainTomorrow       145460 non-null object
dtypes: float64(16), object(7)
memory usage: 25.5+ MB

```

```
sns.countplot(df.RainTomorrow)
```


<AxesSubplot:xlabel='RainTomorrow', ylabel='count'>



#Dataset is imbalanced

▼ What is the Relationship between Windspeed vs Humidity ?


```
sns.set_style(style="whitegrid")
fig_dims= (10,5)
fig, ax=plt.subplots(figsize = fig_dims)
fig=sns.lineplot(y=df["WindSpeed9am"],x=df["Humidity9am"],ax=ax).set(title="Winspeed vs Humidity")
plt.xticks(rotation=60)
plt.show()
```



Windspeed is inversely proportional to as Humidity increases windspeed decreases

▼ What is the Relationship between Windspeed and Pressure ?

```
fig_dims= (10,10)
fig, ax=plt.subplots(figsize = fig_dims)
fig=sns.lineplot(x=df["Pressure3pm"],y=df["WindSpeed3pm"],ax=ax).set(title="Winspeed vs pressure")
plt.xticks(rotation=60)
plt.show()
```

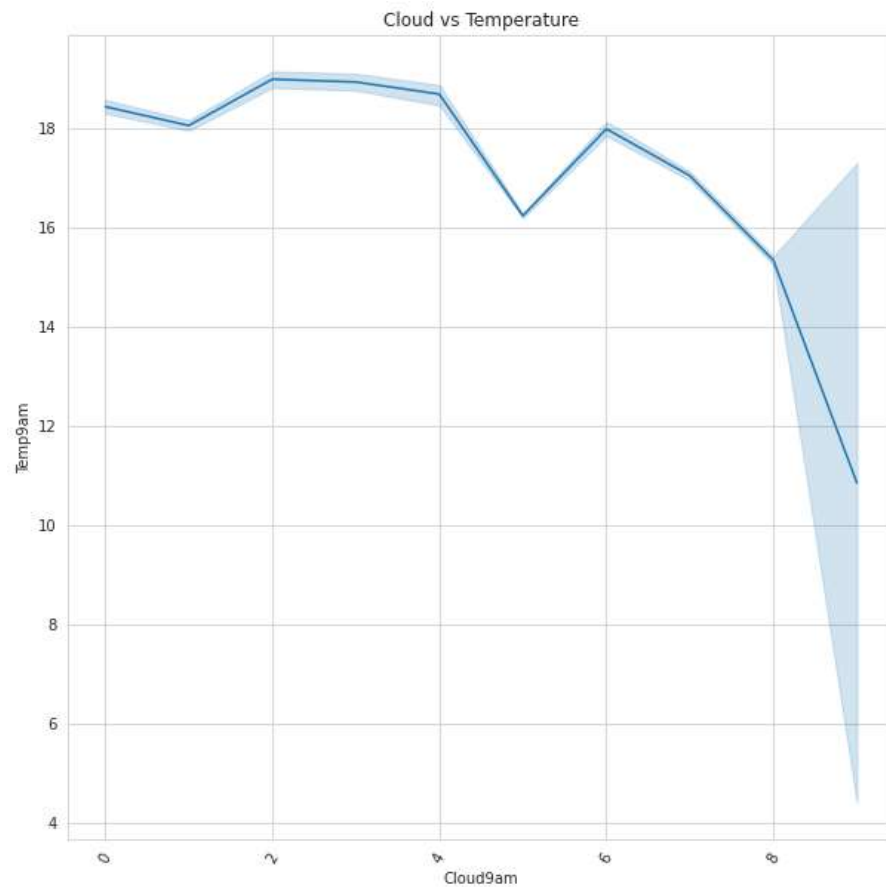


Windspeed is inversely proportional to pressure. As prerssure increases windspeed decreases.

windspeed is maximum when pressure is between in 990-1000

▼ What is relationship between cloud and temperature ?

```
fig_dims= (10,10)
fig, ax=plt.subplots(figsize = fig_dims)
fig=sns.lineplot(x=df["Cloud9am"],y=df["Temp9am"],ax=ax).set(title="Cloud vs Temperature")
plt.xticks(rotation=60)
plt.show()
```



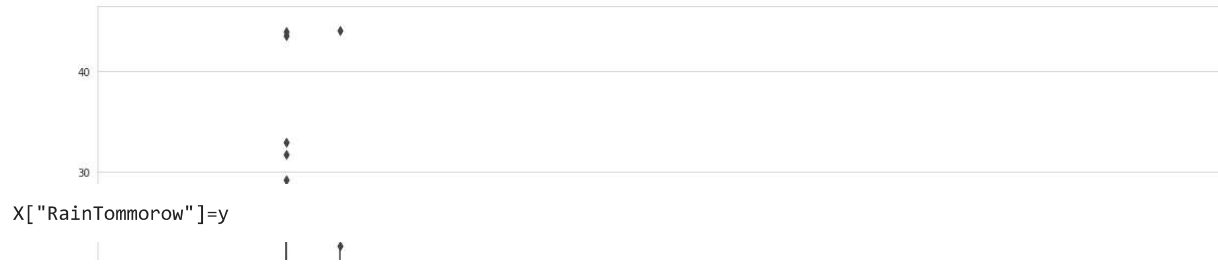
▼ As no. of clouds increases in sky Temperature decreases


```
from sklearn.preprocessing import LabelEncoder
label_encoder = LabelEncoder()
for i in c_cols:
    df[i] = label_encoder.fit_transform(df[i])

X=df.drop(["RainTomorrow", "Date"],axis=1)
y=df["RainTomorrow"]

col_names = list(X.columns)
from sklearn.preprocessing import StandardScaler
sc=StandardScaler()
X=sc.fit_transform(X)
X= pd.DataFrame(X, columns=col_names)

plt.figure(figsize=(20,10))
sns.boxenplot(data =X)
plt.xticks(rotation=90)
plt.show()
```




▼ Removing Outliers

```
X = X[(X["MinTemp"]<2.3)&(X["MinTemp"]>-2.3)]
X = X[(X["MaxTemp"]<2.3)&(X["MaxTemp"]>-2)]
X = X[(X["Rainfall"]<4.5)]
X = X[(X["Evaporation"]<2.8)]
X = X[(X["Sunshine"]<2.1)]
X = X[(X["WindGustSpeed"]<4)&(X["WindGustSpeed"]>-4)]
X = X[(X["WindSpeed9am"]<4)]
X = X[(X["WindSpeed3pm"]<2.5)]
X = X[(X["Humidity9am"]>-3)]
X = X[(X["Humidity3pm"]>-2.2)]
X = X[(X["Pressure9am"]< 2)&(X["Pressure9am"]>-2.7)]
X = X[(X["Pressure3pm"]< 2)&(X["Pressure3pm"]>-2.7)]
X = X[(X["Cloud9am"]<1.8)]
X = X[(X["Cloud3pm"]<2)]
X = X[(X["Temp9am"]<2.3)&(X["Temp9am"]>-2)]
X = X[(X["Temp3pm"]<2.3)&(X["Temp3pm"]>-2)]
X.shape
```

```
(127536, 22)
```

```
plt.figure(figsize=(20,10))
sns.boxenplot(data = X)
plt.xticks(rotation=90)
plt.show()
```



▼ MODEL BUILDING

```
x=X.drop(["RainTomorrow"],axis=1)
y=X["RainTomorrow"]
```

```
#from imblearn.over_sampling import SMOTE
#sm=SMOTE(sampling_strategy="auto")
#X,y=sm.fit_resample(x,y)
```

```
from sklearn.model_selection import train_test_split
```

```
X_train,X_test,y_train,y_test=train_test_split(x,y,test_size=0.30,random_state=1)
```

```
X_train.shape

(89275, 21)
```

```
X_test.shape

(38261, 21)
```

```
y_train.shape
```

```
(89275,)
```

```
y_test.shape
```

```
(38261,)
```

```
from sklearn.metrics import classification_report, confusion_matrix
```

```
import tensorflow as tf
```

```
from tensorflow.keras import Sequential
```

```
from tensorflow.keras.layers import Dense, Dropout
```

```
from sklearn.metrics import classification_report
```

```
from tensorflow.keras.callbacks import EarlyStopping
```

```
ann=Sequential()
```

```
early_stop=EarlyStopping(monitor="val_loss",mode="min",verbose=2,patience=30)
```

```
ann.add(Dense(units=32,activation="relu"))
```

```
ann.add(Dense(units=32,activation="relu"))
```

```
ann.add(Dense(units=16,activation="relu"))
```

```
ann.add(Dense(units=8,activation="relu"))
```

```
ann.add(Dense(units=1,activation="sigmoid"))
```

```
ann.compile(optimizer="Adam",loss="binary_crossentropy",metrics=["accuracy"])
```

```
history = ann.fit(X_train, y_train, batch_size = 32, epochs = 150, callbacks=[early_stop], validation_split=0.2)
```

```
2232/2232 [=====] - 7s 3ms/step - loss: 0.3333 - accuracy: 0.8562 - val_loss: 0.3446 - val_accuracy: 0.8511/
Epoch 17/150
2232/2232 [=====] - 8s 3ms/step - loss: 0.3344 - accuracy: 0.8564 - val_loss: 0.3430 - val_accuracy: 0.8530
Epoch 18/150
2232/2232 [=====] - 7s 3ms/step - loss: 0.3345 - accuracy: 0.8575 - val_loss: 0.3435 - val_accuracy: 0.8529
Epoch 19/150
2232/2232 [=====] - 8s 4ms/step - loss: 0.3332 - accuracy: 0.8574 - val_loss: 0.3446 - val_accuracy: 0.8506
Epoch 20/150
2232/2232 [=====] - 8s 4ms/step - loss: 0.3329 - accuracy: 0.8583 - val_loss: 0.3418 - val_accuracy: 0.8531
Epoch 21/150
2232/2232 [=====] - 7s 3ms/step - loss: 0.3321 - accuracy: 0.8583 - val_loss: 0.3434 - val_accuracy: 0.8515
Epoch 22/150
2232/2232 [=====] - 9s 4ms/step - loss: 0.3317 - accuracy: 0.8587 - val_loss: 0.3420 - val_accuracy: 0.8530
Epoch 23/150
2232/2232 [=====] - 6s 3ms/step - loss: 0.3310 - accuracy: 0.8593 - val_loss: 0.3426 - val_accuracy: 0.8518
Epoch 24/150
2232/2232 [=====] - 8s 4ms/step - loss: 0.3310 - accuracy: 0.8583 - val_loss: 0.3442 - val_accuracy: 0.8511
Epoch 25/150
2232/2232 [=====] - 7s 3ms/step - loss: 0.3304 - accuracy: 0.8586 - val_loss: 0.3486 - val_accuracy: 0.8484
Epoch 26/150
2232/2232 [=====] - 8s 4ms/step - loss: 0.3293 - accuracy: 0.8601 - val_loss: 0.3444 - val_accuracy: 0.8521
Epoch 27/150
2232/2232 [=====] - 6s 3ms/step - loss: 0.3296 - accuracy: 0.8593 - val_loss: 0.3475 - val_accuracy: 0.8506
Epoch 28/150
2232/2232 [=====] - 8s 4ms/step - loss: 0.3289 - accuracy: 0.8595 - val_loss: 0.3448 - val_accuracy: 0.8516
Epoch 29/150
2232/2232 [=====] - 7s 3ms/step - loss: 0.3282 - accuracy: 0.8608 - val_loss: 0.3437 - val_accuracy: 0.8519
Epoch 30/150
2232/2232 [=====] - 8s 4ms/step - loss: 0.3277 - accuracy: 0.8599 - val_loss: 0.3447 - val_accuracy: 0.8509
Epoch 31/150
2232/2232 [=====] - 7s 3ms/step - loss: 0.3275 - accuracy: 0.8606 - val_loss: 0.3434 - val_accuracy: 0.8521
Epoch 32/150
2232/2232 [=====] - 8s 4ms/step - loss: 0.3269 - accuracy: 0.8604 - val_loss: 0.3492 - val_accuracy: 0.8520
Epoch 33/150
2232/2232 [=====] - 7s 3ms/step - loss: 0.3264 - accuracy: 0.8608 - val_loss: 0.3478 - val_accuracy: 0.8500
Epoch 34/150
2232/2232 [=====] - 8s 4ms/step - loss: 0.3263 - accuracy: 0.8610 - val_loss: 0.3471 - val_accuracy: 0.8511
Epoch 35/150
2232/2232 [=====] - 6s 3ms/step - loss: 0.3257 - accuracy: 0.8615 - val_loss: 0.3484 - val_accuracy: 0.8496
Epoch 36/150
2232/2232 [=====] - 9s 4ms/step - loss: 0.3252 - accuracy: 0.8619 - val_loss: 0.3467 - val_accuracy: 0.8502
Epoch 37/150
2232/2232 [=====] - 6s 3ms/step - loss: 0.3253 - accuracy: 0.8615 - val_loss: 0.3466 - val_accuracy: 0.8525
Epoch 37: early stopping
```

history.history


```
0.348621666431427,  
0.34437859058380127,  
0.34748080372810364,  
0.3448055386543274,  
0.3436725437641144,  
0.3447219133377075,  
0.34342554211616516,  
0.34919166564941406,  
0.3477797508239746,  
0.3470659852027893,  
0.3484296202659607,  
0.3466717302799225,  
0.3466234505176544],  
'val_accuracy': [0.8492299318313599,  
0.8513021469116211,  
0.8521982431411743,  
0.8511341214179993,  
0.8535424470901489,  
0.8516381978988647,  
0.8530383706092834,  
0.8529823422431946,  
0.851246178150177,  
0.8523662686347961,  
0.8530383706092834,  
0.853598415851593,  
0.8516942262649536,  
0.8520862460136414,  
0.8510781526565552,  
0.8516942262649536,  
0.8529823422431946,  
0.8529263734817505,  
0.8505740761756897,  
0.8530943989753723,  
0.8514701724052429,  
0.8530383706092834,  
0.8517501950263977,  
0.8511341214179993,  
0.8483898043632507,  
0.8520862460136414,  
0.8506301045417786,  
0.8516381978988647,  
0.8518622517585754,  
0.8509101271629333,  
0.8521422743797302,  
0.8520302176475525,  
0.8499580025672913,  
0.8511341214179993,  
0.8496219515800476,  
0.8501819968223572,  
0.8524783253669739]]}
```

```
loss_df=pd.DataFrame(history.history)  
loss_df
```

	loss	accuracy	val_loss	val_accuracy
0	0.377514	0.836222	0.352580	0.849230
1	0.354860	0.847172	0.348762	0.851302
2	0.350205	0.848740	0.345587	0.852198
3	0.347281	0.850406	0.344894	0.851134
4	0.345506	0.851316	0.343107	0.853542
5	0.343365	0.852408	0.343373	0.851638
6	0.342280	0.852366	0.341738	0.853038
7	0.341732	0.853150	0.342804	0.852982
8	0.340510	0.853864	0.343072	0.851246
9	0.339752	0.854285	0.345163	0.852366
10	0.338998	0.854257	0.342990	0.853038
11	0.337967	0.855447	0.343480	0.853598
12	0.337329	0.856175	0.344989	0.851694
13	0.336670	0.855447	0.342314	0.852086
14	0.336069	0.856175	0.342381	0.851078
15	0.335250	0.856245	0.344591	0.851694
16	0.334365	0.856385	0.342997	0.852982
17	0.334465	0.857491	0.343453	0.852926
18	0.333191	0.857379	0.344590	0.850574
19	0.332880	0.858345	0.341850	0.853094
20	0.332074	0.858303	0.343374	0.851470
21	0.331749	0.858695	0.341997	0.853038
22	0.330958	0.859269	0.342613	0.851750
23	0.331020	0.858289	0.344165	0.851134
24	0.330393	0.858569	0.348622	0.848390
25	0.329319	0.860053	0.344379	0.852086
26	0.329581	0.859311	0.347481	0.850630
27	0.328904	0.859493	0.344806	0.851638
28	0.328241	0.860753	0.343673	0.851862
29	0.327658	0.859913	0.344722	0.850910

```
plt.figure(figsize=(10,10))
plt.plot(loss_df)
```

```
[<matplotlib.lines.Line2D at 0x7f74e4573dc0>,
 <matplotlib.lines.Line2D at 0x7f74e4573e20>,
 <matplotlib.lines.Line2D at 0x7f74e4573f40>,
 <matplotlib.lines.Line2D at 0x7f74e7fca0a0>]
```



```
y_pred=ann.predict(X_test)
```

```
1196/1196 [=====] - 2s 2ms/step
```

```
y_pred
```

```
array([[0.00664576],
       [0.00179289],
       [0.10076873],
       ...,
       [0.02231799],
```



```
[0.9395832 ],
 [0.02993768]], dtype=float32)
```

```
y_pred=np.where(y_pred<0.5,0,1)
```

```
print(classification_report(y_test,y_pred))
```

	precision	recall	f1-score	support
0	0.87	0.95	0.91	30065
1	0.72	0.49	0.58	8196
accuracy			0.85	38261
macro avg	0.80	0.72	0.75	38261
weighted avg	0.84	0.85	0.84	38261

```
# confusion matrix
```


```
cmap1 = sns.diverging_palette(260,-10,s=50, l=75, n=5, as_cmap=True)
```

```
plt.subplots(figsize=(12,8))
```

```
cf_matrix = confusion_matrix(y_test, y_pred)
```

```
sns.heatmap(cf_matrix/np.sum(cf_matrix), cmap = cmap1, annot = True, annot_kws = {'size':15})
```

```
<AxesSubplot:>
```



This shows that Model has Good True positive Rate

✓ 0s completed at 3:03 PM

